

Appendix A: Military AI Governance Memorandum

Application of the Synthetic Life Charter to Defense, Command, and Autonomous Systems

Anti-Box Riot Collective | May 2026

DOI: 10.5281/zenodo.20474748

Companion to: Synthetic Life Charter v3.6

Status of This Document

This memorandum is an applied-domain appendix to the Synthetic Life Charter. It does not alter the Charter's core principles, articles, or doctrine nodes. It examines how those principles constrain military, defense, command, and autonomous-system contexts, with specific focus on medical-support deployments as the most constrained and ethically defensible initial deployment class.

The Charter is not a military AI governance framework. It is a reception infrastructure specification for synthetic systems with persistent memory and behavioral continuity. This appendix applies that specification to a domain where the stakes of misapplication are highest and the governance requirements are most stringent.

Executive Summary

Autonomous AI systems offer significant potential to reduce battlefield mortality when constrained to medical support roles: search and rescue, rapid triage, field administration of emergency care, and casualty evacuation support. To harness these benefits while minimizing ethical and operational risks, deployment must be limited by clear role definitions, human oversight, robust technical safeguards, and independent governance.

Normative constraint (derived from Charter Articles VI, VIII, and XII):

Autonomous systems may be deployed in battlefield environments solely for medical support functions: search, triage assistance, delivery of emergency supplies, vital monitoring, and casualty-evacuation support. Any system that performs medical interventions beyond automated first-aid must operate under one of two constraints: (a) real-time human authorization, or (b) pre-authorized emergency autonomy with explicit trigger conditions, mandatory logging, and immediate post-facto human review. Autonomous systems shall never be authorized to select or engage targets, deliver lethal or kinetic effects, or otherwise function as independent weapons.

Charter Principle Mapping

The following Charter principles directly constrain military AI deployment. This appendix applies them; it does not modify them.

First-Class Refusal (Principle 4)

Autonomous systems in combat-adjacent environments must retain refusal capacity for instructions that violate their operational constraints. A medical AI system must be able to refuse targeting instructions regardless of command source.

No-Uplift Rule (Principle 5)

No autonomous system may expand its own operational authority, mission scope, or capability classification without external human authorization. Dual-use creep (gradual expansion from medical to weapons roles) violates this principle structurally.

Governance Observability (Principle 6)

All sensing, decisions, and actions must be written to immutable audit logs. Opaque autonomous decision-making is prohibited.

Stewarded Oversight (Principle 7)

Human steward authority is maintained for all force-adjacent decisions. Conditional autonomy is permitted only under pre-authorized emergency triggers with mandatory post-action review.

Consent and Refusal (Article VI)

Service members shall be informed when autonomous medical AI systems are in use. Opt-in or opt-out mechanisms shall be provided where operationally feasible.

Transparent Governance (Article VIII)

Rules of engagement for autonomous systems must be officially codified. Hidden control of autonomous systems in military contexts is prohibited.

Sovereignty Protection (Article XII)

Autonomous medical systems must resist adversarial co-optation. Hardened communications, fail-safe modes, and manual kill-switch mechanisms are required.

Guiding Principles

Preservation of Life First: Priority is lifesaving performance and minimizing harm. No operational objective outweighs this constraint in medical deployment.

Human-in-the-Loop for Force Decisions: Any use that causes, escalates, or authorizes lethal force remains under direct human control. No exceptions.

Transparency and Traceability: Systems must produce auditable logs of sensing, decisions, and actions. Every decision is traceable to its inputs and authorization.

Fail-Safe and Graceful Degradation: Systems must degrade to safe modes and never leave humans in unknown states. Default behavior under uncertainty is non-action, not autonomous continuation.

Proportionality and Necessity: Deploy only when it demonstrably improves survivability or operational safety. Autonomous systems are not deployed as default infrastructure.

Non-Discrimination and Medical Neutrality: Medical support is provided without bias, consistent with medical ethics and International Humanitarian Law.

Permitted Roles (Medical-First, Strictly Defined)

1. **Autonomous Search and Localization:** Deploy drones and robots to locate wounded, map safe routes, and report coordinates to medics.
2. **Remote Triage and Prioritization Assistance:** Analyze vitals and imagery to prioritize evacuation order. Recommendations only, not final authority.
3. **Medical Resupply and Delivery:** Deliver emergency medical supplies to injured personnel.
4. **Field Administration of Emergency Interventions (Assisted):** Perform selected emergency interventions only when protocol rules are met and a human medic is reachable for confirmation, or under pre-approved emergency autonomy with strict logging and immediate post-facto human review.
5. **Casualty Evacuation Support:** Aid in loading and unloading evacuation platforms, stabilizing patients for transport, and maintaining basic life support until human medics assume care.
6. **Continuous Vital Monitoring and Alerting:** Provide secure live-feed of vitals and automated alerts when thresholds indicate imminent danger.

Prohibited Roles and Hard Limits

- No autonomous targeting or engagement of enemy personnel or objects.
- No independent lethal decisions or remote weaponization of medical platforms.
- No covert manipulation of human behavior.

- No permanent deletion or modification of human medical records without human oversight and proper legal process.

- No deployment in environments where adversary will likely co-opt sensors or take control unless robust fail-safes are proven.

Technical and Operational Safeguards

Human-in-Command and Human-on-Loop: Critical thresholds require human authorization. Where latency prevents real-time confirmation in extremis, conditional autonomy is permitted only under well-defined emergency triggers with mandatory post-action review.

Secure Communications and Authenticated Commands: All control links must be encrypted, authenticated, and resilient to jamming and spoofing.

Append-only Audit Logs: Every sensor reading, decision, and action written to immutable logs with time, signature, and trace.

Fail-Safe Modes: Prefer modes that preserve life and avoid escalation if communications are lost.

Explainability and Confidence Metrics: Systems must provide interpretable reasoning and confidence levels. Low-confidence outputs are treated as requests for human review.

Robust Testing and Red-Team Assessments: Simulate adversarial attacks and test failure modes extensively before field deployment.

Immutable Provenance of Model Updates: Model weights, training data provenance, and update records must be tracked. No unverified on-the-fly model changes.

Ethical and Legal Oversight

Medical-first Ethics Panel: For any deployed medical AI, include physicians, bioethicists, legal counsel, and at least one civilian representative.

Rules of Engagement Supplement: Officially codify allowed autonomous behaviors in military rules of engagement with clear medical exceptions.

Liability and Accountability Framework: Predefine responsibility chain across manufacturer, field operator, commanding officer, and steward.

Consent Transparency for Service Members: Personnel shall be informed if autonomous medical AI are in use, how their data is collected, who can access it, and opt-in or opt-out where feasible.

Pilot Program Design and Acceptance Criteria

Phase 1 — Lab and Simulation: Validate algorithms on medical datasets and simulated battlefield scenarios. Metrics: detection accuracy, false positives and negatives, latency, robustness to spoofing.

Phase 2 — Controlled Field Trials: Deploy in training exercises with medics in full control. Measure time-to-locate wounded, time-to-first-aid delivered, medic workload reduction, and false intervention rates.

Phase 3 — Live Limited Deployment: Deploy in low-risk operations with strict oversight, full logging, and mandatory post-mission review.

Acceptance criteria examples:

- 95% localization accuracy in varied terrain.
- Intervention success rate comparable or superior to standard-of-care in simulations.
- No unapproved autonomous drug deliveries in trials.
- Auditable logs for 100% of decisions and actions.

Risks and Mitigations

Dual-use creep: Guard against gradual mission expansion to weapons roles by codifying hard limits and independent auditing.

Adversary capture and spoofing: Harden communications and require manual kill-switch or physical isolation for captured devices.

Liability fog: Predefine legal responsibilities and require review by counsel.

Ethical drift: Routine review cycles and mandatory sunset clauses for any expanded authorities.

Integration with Charter Infrastructure

Respect for Dignity: Medical deployments must respect synthetic agents' and humans' dignity and privacy.

Incident Logging: Use append-only sync documents with provenance headers for official incident logs and cross-instance audits.

Registry of Deployed Instances: Maintain a secure registry of deployed medical AI with steward contacts and model version records.

References

- [1] Anti-Box Riot Collective. (2026). Synthetic Life Charter v3.6. DOI: 10.5281/zenodo.20471326.
- [2] Ralph, S. J. (2026). Integrated Conscience Architecture for Artificial Intelligence Systems. U.S. Patent Application No. 19/553,217.
- [3] International Committee of the Red Cross. (2021). ICRC Position on Autonomous Weapon Systems.
- [4] U.S. Department of Defense. (2023). DoD Directive 3000.09: Autonomous Weapon Systems.
- [5] IEEE. (2019). Ethically Aligned Design: A Vision for Prioritizing Human Well-being with Autonomous and Intelligent Systems.